

APPENDIX for MIND: A Coreset Selection for Continual Crowd Trajectory Prediction

I. EXPERIMENT CONFIGURATION

On Reducible Loss, the subset model is trained for 100 epochs; under our combined MIND method, the equivalent subset model requires only 10 epochs. On BCSR, we run 10 outer and 10 inner iterations; under MIND, we use only 5 outer and 5 inner iterations. MIND therefore requires substantially fewer training epochs and iterations per task than either coreset selection method used in isolation, reflecting that combining the three scores allows the coreset to reach a comparably useful state more quickly than relying on a single scoring signal alone. For Experience Learning, we use a uniform distribution to select the subset from the data source.

CPTL is trained using its unmodified training procedure, since it relies on its own generative replay model rather than our coreset selection pipeline. However, to have an efficient experiment since CPTL requires huge GPU memory, we trained CPTL on the dynamic Cluster dataset.

For trajectory prediction, we use 24 observed steps and a 32-step future horizon, longer than the standard 8-observed/12-future convention used in most prior work (e.g., ETH-UCY-based evaluation), in order to more clearly test long-horizon prediction performance under continual learning. We do not extend this horizon further, since a larger observation-and-prediction window would exceed the length of some trajectories in our dataset, leaving insufficient data to support training at that horizon.

To confirm the robustness of the results against random effects, we repeated the evaluation process 5 times for all experiments, using 5 different random seeds (1 through 5), and report results as mean $\pm$ standard deviation. On the trade-off calculation we did the normalization for whole methods including the ablation on Table V

We used different machine on runpod.io which share the same GPU which is 5090 but vary on CPU to conduct each experiment to make the experiment become parallel.

II. ADDITIONAL EVALUATION METRICS

This is additional evaluation metrics to calculate total forgetting: **Final Forgetting** measures the final model’s forgetting across all previously seen scenes. For each scene i , we compute the difference between the final model’s error on scene i ’s data and the error recorded when scene i was first learned. The last scene is excluded, since it has only been trained on once.

$$\text{FinalForgetting} = \frac{1}{N-1} \sum_{i=1}^{N-1} (e_{i,N} - e_{i,i}) \quad (1)$$

where N is the number of tasks (scenes), and $e_{i,j}$ is the ADE on scene i ’s data using the model checkpoint after training on scene j .

III. DETAILED RESULT

Tables I–IV show that our method outperforms other continual learning methods, particularly in its ability to preserve previously learned knowledge.

A. Fluctuating Performance of Replay-Agnostic Methods

Comparing the two replay-agnostic baselines, Independent Learning and Zero-Shot, we observe fluctuating relative Total Forgetting performance between them. Independent-Full requires more trajectory data to train, since its data is not reduced by the preprocessing dynamic clustering (see Table ??) and it uses the full dataset for every task. Meanwhile, Zero-Shot, which applies our pipeline’s filtering directly, runs with less trajectory data. On Plasticity, however, the pattern is more consistent than for forgetting: Independent CL has the worst Plasticity of the three methods on every task ordering, while Zero-Shot CL matches or outperforms it in every case and remains close to Independent-Full. This shows that filtering does not hurt plasticity compared to relying only on the base model’s own generalization ability. In fact, filtering gives better plasticity than simply reducing the data through clustering alone.

B. Performance on the Full Data Source

On the full data source, Independent-Full achieves the best Total Forgetting in three of the four task orderings, reflecting that it relies entirely on the base model’s own generalization ability, with no replay mechanism. In particular, this result holds across all orderings for the SocialVAE backbone, which uses a limited neighbourhood representation, outperforming the MART backbone, which models the full neighbourhood jointly. Since Independent-Full has no replay data at all, we use its SocialVAE variant as our reference point for plasticity performance. This reference, however, comes at a computational cost: Independent-Full requires substantially more computation than REL and BCSR as applied within our pipeline, which retain only 25% of the data at each stage while still operating on the full (unclustered) trajectory representation. Applying our pipeline’s filtering to REL and BCSR reduces their Total Forgetting performance relative to Independent-Full (BCSR achieves the best performance only in Table IV); on plasticity, however, both outperform Independent-Full in three of the four orderings.

TABLE I: Crowd-Traj(DM-DE-V-D1-P). Continual trajectory prediction performance under the first order. Lower values indicate better performance for all metrics. The table shows that our proposed algorithm has the best performance on forgetting scores and also for the stability-plasticity tradeoff.

Method	DS	Total Forgetting	Avg. Error	Last Forgetting	Plasticity	Tradeoff
Independent	Full	0.0059 $\pm$ 1.6e-05	0.0633 $\pm$ 1.1e-05	0.0152 $\pm$ 4.4e-05	0.0486 $\pm$ 1.0e-05	0.1899
Ind.(MART)	Full	0.0101 $\pm$ 2.3e-05	0.0840 $\pm$ 1.5e-05	0.0289 $\pm$ 6.0e-05	0.0533 $\pm$ 1.4e-05	0.4029
REL	Full	0.0064 $\pm$ 1.1e-05	0.0629 $\pm$ 5.8e-06	0.0159 $\pm$ 4.0e-05	0.0470 $\pm$ 1.8e-05	0.2022
BCSR	Full	0.0072 $\pm$ 1.4e-05	0.0635 $\pm$ 2.0e-05	0.0178 $\pm$ 3.5e-05	0.0470 $\pm$ 2.1e-05	0.2374
CPTL(RNN)	CL	0.0248 $\pm$ 2.1e-05	0.2967 $\pm$ 4.6e-05	0.0683 $\pm$ 2.8e-05	0.2156 $\pm$ 4.0e-05	2.0000
Independent	CL	0.0053 $\pm$ 1.9e-05	0.0661 $\pm$ 1.9e-05	0.0142 $\pm$ 4.2e-05	0.0509 $\pm$ 1.5e-05	0.1799
Zero Shot	CL	0.0113 $\pm$ 1.0e-05	0.0693 $\pm$ 2.3e-05	0.0345 $\pm$ 1.4e-05	0.0488 $\pm$ 2.0e-05	0.4257
Experience	CL	0.0047 $\pm$ 1.6e-05	0.0649 $\pm$ 1.6e-05	0.0143 $\pm$ 3.5e-05	0.0499 $\pm$ 1.2e-05	0.1469
BCSR	CL	0.0052 $\pm$ 2.0e-05	0.0627 $\pm$ 1.5e-05	0.0155 $\pm$ 2.9e-05	0.0499 $\pm$ 1.2e-05	0.1641
REL	CL	0.0041 $\pm$ 3.1e-05	0.0621 $\pm$ 2.6e-05	0.0111 $\pm$ 7.8e-05	0.0499 $\pm$ 2.1e-05	0.1195
MIND(ours)	CL	0.0017 $\pm$ 9.2e-05	0.0583 $\pm$ 2.2e-05	0.0055 $\pm$ 2.6e-05	0.0508 $\pm$ 2.7e-05	0.0225

TABLE II: Crowd-Traj(D1-V-DM-P-DE). Continual trajectory prediction performance under the second order. Lower values indicate better performance for all metrics. The table shows that our proposed algorithm has the best performance on forgetting scores and also for the stability-plasticity tradeoff.

Method	DS	Total Forgetting	Avg. Error	Last Forgetting	Plasticity	Tradeoff
Independent	Full	0.0138 $\pm$ 1.8e-05	0.0585 $\pm$ 3.5e-05	0.0272 $\pm$ 2.2e-05	0.0486 $\pm$ 3.8e-05	0.2560
Ind.(MART)	Full	0.0224 $\pm$ 4.6e-05	0.0736 $\pm$ 2.4e-05	0.0442 $\pm$ 8.8e-05	0.0533 $\pm$ 1.2e-05	0.5135
REL	Full	0.0140 $\pm$ 2.4e-05	0.0568 $\pm$ 2.0e-05	0.0273 $\pm$ 7.3e-05	0.0464 $\pm$ 9.8e-06	0.2475
BCSR	Full	0.0151 $\pm$ 3.8e-05	0.0576 $\pm$ 2.0e-05	0.0290 $\pm$ 7.3e-05	0.0460 $\pm$ 9.8e-06	0.2745
CPTL(RNN)	CL	0.0425 $\pm$ 1.0e-04	0.2422 $\pm$ 2.2e-05	0.0853 $\pm$ 1.6e-04	0.2089 $\pm$ 4.6e-05	2.0000
Independent	CL	0.0138 $\pm$ 4.3e-05	0.0597 $\pm$ 9.1e-06	0.0257 $\pm$ 6.5e-05	0.0509 $\pm$ 1.9e-05	0.2702
Zero Shot	CL	0.0152 $\pm$ 4.7e-05	0.0578 $\pm$ 1.0e-05	0.0284 $\pm$ 7.2e-05	0.0482 $\pm$ 2.1e-05	0.2906
Experience	CL	0.0115 $\pm$ 2.6e-05	0.0560 $\pm$ 1.6e-05	0.0222 $\pm$ 5.0e-05	0.0482 $\pm$ 2.8e-05	0.1943
BCSR	CL	0.0165 $\pm$ 1.9e-05	0.0590 $\pm$ 2.4e-05	0.0319 $\pm$ 4.1e-05	0.0486 $\pm$ 2.3e-05	0.3275
REL	CL	0.0120 $\pm$ 2.4e-05	0.0564 $\pm$ 3.3e-05	0.0236 $\pm$ 3.2e-05	0.0486 $\pm$ 2.3e-05	0.2106
MIND(ours)	CL	0.0057 $\pm$ 2.2e-05	0.0541 $\pm$ 9.4e-06	0.0118 $\pm$ 3.6e-05	0.0510 $\pm$ 2.5e-05	0.0573

TABLE III: Julich&MOT21(b2-d7-g3-M2). Continual trajectory prediction performance under this task order. Lower values indicate better performance for all metrics. The table shows that our proposed algorithm has the best performance on forgetting scores and also for the stability-plasticity tradeoff.

Method	DS	Total Forgetting	Avg. Error	Last Forgetting	Plasticity	Tradeoff
Independent	Full	0.6936 $\pm$ 1.8e-03	3.6041 $\pm$ 1.5e-03	1.1032 $\pm$ 2.3e-03	2.9182 $\pm$ 6.2e-04	0.2557
Ind.(MART)	Full	1.0794 $\pm$ 2.9e-03	4.2810 $\pm$ 2.4e-03	1.6607 $\pm$ 3.6e-03	3.1448 $\pm$ 9.7e-04	0.4420
REL	Full	0.7701 $\pm$ 2.9e-03	3.4245 $\pm$ 1.7e-03	1.2234 $\pm$ 3.7e-03	2.7734 $\pm$ 7.0e-04	0.2460
BCSR	Full	0.7718 $\pm$ 2.0e-03	3.4315 $\pm$ 9.4e-04	1.2172 $\pm$ 3.1e-03	2.8045 $\pm$ 1.6e-03	0.2542
CPTL(RNN)	CL	2.9988 $\pm$ 1.9e-03	9.2254 $\pm$ 2.2e-03	4.6573 $\pm$ 2.3e-03	6.8379 $\pm$ 2.6e-03	2.0000
Independent	CL	2.3074 $\pm$ 1.6e-03	4.4713 $\pm$ 1.8e-03	3.5712 $\pm$ 1.9e-03	3.1467 $\pm$ 2.1e-03	0.8579
Zero Shot	CL	0.4962 $\pm$ 1.7e-03	3.4972 $\pm$ 1.3e-03	0.7735 $\pm$ 2.8e-03	2.9825 $\pm$ 1.4e-03	0.2047
Experience	CL	0.5050 $\pm$ 5.9e-04	3.5299 $\pm$ 1.6e-03	0.8260 $\pm$ 1.8e-03	2.9897 $\pm$ 2.3e-03	0.2095
BCSR	CL	0.4304 $\pm$ 1.7e-03	3.4646 $\pm$ 1.9e-03	0.6917 $\pm$ 3.2e-03	2.9585 $\pm$ 2.7e-03	0.1766
REL	CL	0.5019 $\pm$ 3.7e-03	3.5282 $\pm$ 1.2e-03	0.8231 $\pm$ 5.4e-03	2.9899 $\pm$ 1.8e-03	0.2085
MIND(ours)	CL	0.1070 $\pm$ 1.5e-03	3.2102 $\pm$ 7.3e-04	0.2051 $\pm$ 3.1e-03	2.9487 $\pm$ 2.0e-03	0.0648

Comparing REL and BCSR directly on the full data source, both Total Forgetting and Plasticity fluctuate between the two across orderings: REL achieves better Total Forgetting in three of the four orderings, while on Plasticity,

BCSR is only marginally ahead in the first two orderings before REL takes a clearer lead in the third and fourth. This indicates that their relative performance on the full dataset depends on task ordering rather than one method

TABLE IV: Jülich&MOT21(g3-d7-M2-b2). Continual trajectory prediction performance under this task order. Lower values indicate better performance for all metrics. The table shows that our proposed algorithm has the best performance on forgetting scores and also for the stability-plasticity tradeoff.

Method	DS	Total Forgetting	Avg. Error	Last Forgetting	Plasticity	Tradeoff
Independent	Full	$0.3807 \pm 2.1\text{e-}03$	$3.6602 \pm 1.4\text{e-}03$	$0.4912 \pm 4.0\text{e-}03$	$2.9164 \pm 1.6\text{e-}03$	0.1754
Ind.(MART)	Full	$0.5778 \pm 3.2\text{e-}03$	$4.4536 \pm 2.0\text{e-}03$	$1.0120 \pm 6.0\text{e-}03$	$3.1448 \pm 2.3\text{e-}03$	0.3277
REL	Full	$0.3888 \pm 2.6\text{e-}03$	$3.7480 \pm 2.1\text{e-}03$	$0.4802 \pm 5.0\text{e-}03$	$3.0019 \pm 2.1\text{e-}03$	0.1952
BCSR	Full	$0.2714 \pm 1.4\text{e-}03$	$3.6713 \pm 1.3\text{e-}03$	$0.3399 \pm 3.0\text{e-}03$	$3.0767 \pm 1.7\text{e-}03$	0.1421
CPTL(RNN)	CL	$1.8419 \pm 9.2\text{e-}03$	$12.3314 \pm 1.0\text{e-}02$	$3.3448 \pm 2.8\text{e-}02$	$8.5103 \pm 1.2\text{e-}02$	2.0000
Independent	CL	$0.2873 \pm 1.2\text{e-}03$	$4.3982 \pm 1.4\text{e-}03$	$0.4455 \pm 2.6\text{e-}03$	$3.1464 \pm 2.0\text{e-}03$	0.1635
Zero Shot	CL	$0.3431 \pm 1.5\text{e-}03$	$3.7332 \pm 1.7\text{e-}03$	$0.4978 \pm 4.7\text{e-}03$	$3.0722 \pm 2.0\text{e-}03$	0.1819
Experience	CL	$0.3709 \pm 1.1\text{e-}03$	$3.5557 \pm 1.5\text{e-}03$	$0.5533 \pm 3.3\text{e-}03$	$2.9000 \pm 1.6\text{e-}03$	0.1669
BCSR	CL	$0.3185 \pm 1.1\text{e-}03$	$3.5690 \pm 1.6\text{e-}03$	$0.4834 \pm 3.0\text{e-}03$	$2.9333 \pm 1.9\text{e-}03$	0.1432
REL	CL	$0.3709 \pm 2.7\text{e-}03$	$3.5560 \pm 1.3\text{e-}03$	$0.5531 \pm 6.8\text{e-}03$	$2.8999 \pm 2.5\text{e-}03$	0.1669
MIND(ours)	CL	$0.0957 \pm 6.6\text{e-}04$	$3.5198 \pm 9.1\text{e-}04$	$0.1428 \pm 1.4\text{e-}03$	$3.1160 \pm 1.2\text{e-}03$	0.0496

being consistently superior.

C. Performance on the Dynamic Cluster Data Source

This data source reflects the pre-processing step in our pipeline, in which raw trajectory data is reduced and represented via cluster centroids, resulting in substantially fewer trajectories than the full data source.

Comparing Independent-Cluster against Independent-Full, plasticity degrades on all four task orderings, with the gap remaining substantial even on the fourth ordering (2.9164 for Independent-Full versus 3.1464 for Independent-Cluster). Given the substantial reduction in trajectory volume achieved by clustering, this consistent but modest plasticity cost supports dynamic clustering as an effective efficiency step with a favourable overall trade-off.

We evaluate CPTL under the same continual learning setting, though its generative replay mechanism requires substantially more memory than our pipeline’s filtering approach, since it must both train a generative model and collect the resulting synthetic training data. CPTL shows the worst plasticity of any method tested, with a substantial gap relative to both the SocialVAE and MART backbones, resulting in the worst Trade-off score across all task orderings.

Zero-Shot is, in principle, the most computationally efficient method evaluated, since it involves no replay computation at all, applying only our pipeline’s base filtering step. Compared to the Independent-Full baseline, Zero-Shot achieves better Total Forgetting on the Jülich/MOT21 orderings (Tables III and IV), indicating that the trade-off between prediction error and dataset reduction is worthwhile even without any replay mechanism, although it does not achieve the best Total Forgetting overall on those orderings once BCSR, Independent-Cluster, and MIND are considered. On the dynamic cluster data source specifically, Zero-Shot also achieves the best plasticity in Tables I and II, though this advantage does not hold when compared against methods using the full data source.

MIND achieves the best Total Forgetting of any method, on every data source and task ordering tested, by a substantial

margin (0.0017 on Table I, 0.0057 on Table II, 0.1070 on Table III, and 0.0957 on Table IV). This comes at the expected trade-off in plasticity on three of the four orderings; the exception is Table III, on the dynamic cluster data source, where MIND achieves the best plasticity as well, alongside the best average error and Last Forgetting, indicating strong overall error control and forgetting reduction with no corresponding cost to adaptability on that ordering.

Comparing MIND’s average error on the dynamic cluster data source (CL) against methods using the full data source, MIND achieves lower average error than every full-source method on all four task orderings. This indicates that our coreset-based replay mechanism is sufficient to control prediction error even without access to the larger, more diverse full trajectory set, suggesting that explicit, criteria-driven replay is a more reliable source of error control than simply training on more unfiltered data.

IV. CORESET COMPOSITION ANALYSIS

Figure 1 compares the loss score distribution of the full dataset (blue) against the coreset (red) for each method. In the upper graph, REL’s distribution retains more samples at high loss scores, visible as a longer tail extending to higher values, consistent with REL’s selection criterion placing greater emphasis on high reducible-loss samples. This tail converges gradually rather than dropping off sharply. By contrast, the BCSR and MIND coresets retain the fewest samples at high loss scores, shown by the clear gap between the tail of the full dataset’s distribution and that of each coreset.

Figure 2 provides further detail, showing the same loss score distributions with the y-axis scaled by a factor of 10 in each panel. As in Figure 1, most of the distribution is concentrated at low loss scores, indicating that the model handles the majority of tasks well. MIND’s coreset shows the lowest overall loss score distribution of the three methods, with its highest value reaching approximately 0.4 out of 1.4 compared to all data; REL’s most extreme sample reaches a loss score of approximately 4 out of 4 compared to all data,

and BCSR’s most extreme sample reaches approximately 1.4 out of 8 compared to all data. This indicates that MIND’s coreset construction is comparatively insensitive to extreme loss values, resulting in a coreset that does not retain the most extreme high-loss outliers to the same extent as REL or BCSR.

The lower graph in Figure 1 maps each coreset’s internal composition across the continuous range of cluster-member counts observed, expressed as a percentage of each method’s own coreset. Since MIND, REL, and BCSR are constructed with different target coreset sizes, these percentages describe the relative shape of each method’s selection, rather than absolute sample counts, and should not be compared directly across methods as raw quantities.

Within each method’s own coreset, MIND shows the strongest skew toward low cluster-membership samples: 89.6% of MIND’s coreset consists of samples with only 1 cluster member, compared to 87.946% for REL and 87.55% for BCSR. At the opposite end of the range, MIND allocates the smallest proportion of its coreset to samples with 6 cluster members (0.074%), compared to REL (0.134%) and BCSR (0.178%), the largest of the three.

This distributional skew indicates that, relative to their respective coreset sizes, MIND is proportionally less inclined than REL or BCSR to select high-cluster-membership samples, and proportionally more inclined to select low-membership samples. A high cluster-membership count alone does not guarantee selection if a sample scores poorly on Influence or Difficulty, and conversely, a sample need not have many cluster members to be selected if it scores well on the other two criteria. This is also consistent with the loss-distribution results above: because MIND requires a sample to be evaluated well across all three criteria jointly, rather than any single criterion in isolation, it avoids over-selecting samples purely based on high loss (as REL tends to) or high cluster membership alone, resulting in the more balanced, lower-variance coreset composition observed in both figures.

V. MIND PARAMETER PAIR ROLES

The ablation results across the four task sequences demonstrate that the contribution of Membership, Influence, and Difficulty varies depending on task ordering, while the complete MIND configuration provides a generally robust balance between forgetting and plasticity. MIND achieves the lowest Trade-off in two of the four task sequences: the first CrowdTraj sequence (DM-DA-V-D1-P, 0.0225) and the first Jülich&MOT21 sequence (b2-d7-g3-M2, 0.0648). This provides evidence that jointly incorporating all three criteria is particularly effective for these task orderings. However, MIND does not consistently achieve the best performance across all task orders: on the second CrowdTraj sequence (D1-V-DM-P-DE), IND achieves the lowest Trade-off (0.0333), and on the second Jülich&MOT21 sequence (g3-d7-M2-b2), MD achieves the lowest Trade-off (0.0146). This variation indicates that the contribution of each criterion depends on the characteristics and ordering of the tasks.

From a computational perspective, the ablation results also suggest that removing computationally expensive criteria can offer a favourable performance-efficiency trade-off in certain settings. Among the reduced configurations (IND, MD, MIN), IND achieves the lowest Trade-off on both CrowdTraj sequences (0.0741 and 0.0333), while MD achieves the lowest Trade-off on both Jülich&MOT21 sequences (0.0823 and 0.0146) — the latter also being the overall lowest Trade-off across all four configurations on that sequence. This suggests that the relative usefulness of Influence versus Membership depends on the dataset: on CrowdTraj, retaining Influence and Difficulty (IND) is more effective than retaining Membership and Difficulty (MD), whereas the reverse holds on Jülich&MOT21. Because Influence requires additional computation during sample evaluation, MD may represent a more computationally efficient alternative to the full MIND configuration since it does not need to compute the influence score.




Fig. 1: The upper graph shows a comparison of core-set selection distribution between BCSR, REL, and MIND(ours) methods based on the last loss score (scaled 10 times) from main training. The lower graph shows a comparison based on the cluster score gained from pre-processing with dynamic clustering.




Fig. 2: The right graph shows the coreset distribution on one dimension of loss score; the left graph shows the whole dataset loss score. These images are sampled for 120b2 datasets from the Julich dataset for MIND (our proposed), BCSR, and REL on the CL data source.

TABLE V: Performance of different parameter configurations of MIND across dataset sequences. MIND uses all parameters, while IND uses Influence and Difficulty, MD uses Membership and Difficulty, and MIN uses Membership and Influence. Reported as mean $\pm$ standard deviation over 5 training runs.

(a) First task sequence: CrowdTraj: DM-DA-V-D1-P

Method	Total Forgetting	Avg. Accuracy	Last Forgetting	Plasticity	Trade-off
MIND	0.0017 $\pm$ 9.2e-06	0.0583 $\pm$ 2.2e-05	0.0055 $\pm$ 2.6e-05	0.0508 $\pm$ 2.7e-05	0.0225
IND	0.0030 $\pm$ 2.4e-05	0.0587 $\pm$ 1.3e-05	0.0092 $\pm$ 4.8e-05	0.0497 $\pm$ 3.2e-05	0.0741
MD	0.0031 $\pm$ 2.8e-05	0.0582 $\pm$ 2.1e-05	0.0091 $\pm$ 5.8e-05	0.0498 $\pm$ 3.0e-05	0.0792
MIN	0.0075 $\pm$ 2.8e-05	0.0677 $\pm$ 2.2e-05	0.0226 $\pm$ 6.5e-05	0.0493 $\pm$ 3.0e-05	0.2625

(b) Second task sequence: CrowdTraj: D1-V-DM-P-DE

Method	Total Forgetting	Avg. Accuracy	Last Forgetting	Plasticity	Trade-off
MIND	0.0057 $\pm$ 2.2e-05	0.0541 $\pm$ 9.4e-06	0.0118 $\pm$ 3.6e-05	0.0510 $\pm$ 2.5e-05	0.0573
IND	0.0047 $\pm$ 2.3e-05	0.0541 $\pm$ 1.8e-05	0.0101 $\pm$ 5.3e-05	0.0514 $\pm$ 2.2e-05	0.0333
MD	0.0078 $\pm$ 2.3e-05	0.0602 $\pm$ 1.3e-05	0.0162 $\pm$ 5.9e-05	0.0524 $\pm$ 1.3e-05	0.1202
MIN	0.0137 $\pm$ 3.8e-05	0.0585 $\pm$ 1.8e-05	0.0257 $\pm$ 7.4e-05	0.0487 $\pm$ 2.0e-05	0.2531

(c) Third task sequence: Julich&MOT21: b2-d7-g3-M2

Method	Total Forgetting	Avg. Accuracy	Last Forgetting	Plasticity	Trade-off
MIND	0.1070 $\pm$ 1.5e-03	3.2102 $\pm$ 7.3e-04	0.2051 $\pm$ 3.1e-03	2.9487 $\pm$ 2.0e-03	0.0648
IND	0.0432 $\pm$ 9.6e-04	3.4998 $\pm$ 1.2e-03	0.1167 $\pm$ 2.8e-03	3.1906 $\pm$ 1.1e-03	0.1027
MD	0.0677 $\pm$ 1.6e-03	3.3507 $\pm$ 1.2e-03	0.1184 $\pm$ 4.1e-03	3.0743 $\pm$ 2.3e-03	0.0823
MIN	0.4622 $\pm$ 2.6e-03	3.6798 $\pm$ 7.2e-04	0.7220 $\pm$ 5.0e-03	3.1244 $\pm$ 1.7e-03	0.2281

(d) Fourth task sequence: Julich&MOT21: g3-d7-M2-b2

Method	Total Forgetting	Avg. Accuracy	Last Forgetting	Plasticity	Trade-off
MIND	0.0957 $\pm$ 6.6e-04	3.5198 $\pm$ 9.1e-04	0.1428 $\pm$ 1.4e-03	3.1160 $\pm$ 1.2e-03	0.0496
IND	0.1021 $\pm$ 9.5e-04	3.3875 $\pm$ 1.6e-03	0.2009 $\pm$ 3.0e-03	2.9593 $\pm$ 1.6e-03	0.0252
MD	0.0762 $\pm$ 1.4e-03	3.3438 $\pm$ 8.5e-04	0.1522 $\pm$ 3.4e-03	2.9817 $\pm$ 1.1e-03	0.0146
MIN	0.3220 $\pm$ 9.5e-04	3.5599 $\pm$ 9.1e-04	0.4812 $\pm$ 2.7e-03	2.9078 $\pm$ 1.6e-03	0.1406